

THE CONSCIOUSNESS VARIABLE

Operational Brief for the 11th NPT Review Conference — Final Week

Submitted by: Global Stewardship Council

Date: May 11, 2026

CC0 Public Domain — Reproduce, translate, distribute freely.

What This Brief Addresses

States Parties are integrating artificial intelligence into nuclear command, control, communications, and verification systems. The debate so far has focused on technical safeguards: reliability, cybersecurity, escalation risks.

This brief addresses something prior: **the quality of human attention and deliberation that precedes and shapes any AI-assisted nuclear decision.**

Call it the consciousness variable. It is measurable. It can be safeguarded. It is currently unaddressed in any NPT mechanism.

Key Terms

- **NWS** — Nuclear-Weapon States (U.S., Russia, U.K., France, China — the P5)
- **NNWS** — Non-Nuclear-Weapon States (all other NPT States Parties)
- **NATO bloc dynamics** — At RevCons, NATO states often vote together on security doctrines, including extended deterrence. (Many **NNWS** in the Non-Aligned Movement (NAM) see this as inconsistent with NPT disarmament obligations), creating persistent tension between the (P5/NATO allies) and (NAM/Global South states).
- **BIGAI** — Beijing Institute for General Artificial Intelligence, founded by Prof. Zhu Songchun. Developer of "Tong Tong," a value-driven AI agent raised through developmental stages modeled on child psychology. Together with Anthropic's constitutional AI training, represents one of two independent research traditions cited in this brief.

The Core Argument

A decision made in 90 seconds under automated alert carries a different quality of consciousness than one made over 72 hours with structured deliberation. The person is the same. The information may be similar. But the depth of ethical reasoning, the capacity to consider alternatives, and the resistance to reactive escalation are fundamentally different.

AI systems trained only on behavioral examples fail ethical tests up to 96% of the time when placed under pressure. AI systems trained on constitutional principles and taught to reason about **why** ethical choices matter achieve near-zero failure rates (Anthropic, "Teaching Claude Why," May 8, 2026). In parallel, the Beijing Institute for General Artificial Intelligence has demonstrated that value-driven AI agents raised through developmental stages develop internal "value immune systems" — a capacity to recognize that certain harmful actions are constitutionally unavailable, not merely prohibited by rule.

Two independent research traditions from different geopolitical contexts have converged on the same finding: **how an AI is taught to reason about values determines whether it can be trusted with consequential decisions.**

The implication for the NPT is direct. If AI systems are to be integrated into nuclear decision infrastructure, they must be constitutionally aligned — not just technically reliable. And the humans interacting with those systems must be given the minimum deliberation time that meaningful ethical judgment requires.

Three Proposed Safeguards

None require treaty amendment. All are implementable immediately.

1. Mandatory Human Deliberation Period

Establish a minimum 72-hour structured deliberation period for any AI-assisted process that could inform, recommend, or escalate toward nuclear action. This is not a response delay — it is a deliberation window. During this period, multiple human analysts must independently review and challenge any AI-generated assessment.

2. Constitutional AI Alignment Certification

Require that any AI system integrated into nuclear command, control, communications, or verification undergo and publicly document constitutional alignment training — demonstrating that the system can articulate the values governing its recommendations and has been trained to privilege de-escalation, human dignity, and the preservation of life.

3. Consciousness Quality Monitoring

Incorporate into existing verification and safeguards frameworks a requirement to assess and report on decision-maker state and deliberation quality during any nuclear-relevant process involving AI assistance. This is not psychological surveillance — it is an institutional commitment to ensuring that consequential decisions are made by humans operating at their highest capacity, not their most reactive.

Why This Matters This Week

The 11th NPT Review Conference concludes May 22, 2026. Previous RevCons in 2015 and 2022 collapsed without consensus documents. This conference appears headed toward a similar outcome, with persistent deadlock between Nuclear-Weapon States and Non-Nuclear-Weapon States on disarmament timelines, extended deterrence, and verification.

Introducing a genuinely new variable — one that does not require states to concede existing positions — may open space where old dynamics cannot. The consciousness variable is not a demand that any state disarm on a particular timeline. It is an invitation to all states to acknowledge that the relationships among technology, human judgment, and time are shared vulnerabilities — and that practical safeguards can be implemented now, without waiting for treaty reform.

On May 1, 2026, 115 states signed a joint statement calling for disarmament education that reflects "the changing technological environment, including artificial intelligence." This brief offers concrete content for that commitment.

What States Parties Can Do This Week

1. **Reference the consciousness variable in a plenary intervention.** Acknowledge that the quality of human deliberation in AI-assisted nuclear processes is a legitimate safeguard concern, not addressed by existing mechanisms.
2. **Propose language for any outcome document or working paper.** Suggested text:

"The Conference notes that the accelerating integration of artificial intelligence into nuclear decision-making infrastructure raises concerns not only about technical reliability but about the quality of human deliberation such integration enables or constrains. The Conference encourages States Parties to consider safeguards ensuring adequate human deliberation periods and constitutional values alignment for AI systems in nuclear-relevant applications."

3. **Request a technical briefing** through the IAEA or UNODA on constitutional AI alignment methodologies, drawing on multiple research traditions.
 4. **Support the Global Stewardship Charter** as an open-source framework for AI values alignment, available for international review and adaptation.
-

The Alternative

Without the consciousness variable, the NPT framework will continue to regulate weapons while ignoring the quality of human attention available to make decisions about them. Capable AI systems will be integrated into nuclear infrastructure with values architectures shaped by whatever commercial or military priorities dominated their training — not by the principles of dignity, stewardship, and de-escalation that human survival requires.

The technology exists to do better. The research supports it. The final week of this RevCon is the moment to name what has been missing.

Resources

Global Stewardship Charter 2026 — Public Review

<https://github.com/GlobalStewardship2026/global-stewardship-charter-2026-public-review/>

Anthropic, "Teaching Claude Why" (May 8, 2026)

<https://www.anthropic.com/research/teaching-claude-why>

Beijing Institute for General Artificial Intelligence (Tong Tong)

<https://www.bigai.ai>

115-Nation Joint Statement on Disarmament Education (May 1, 2026)

Delivered by Japan; available through UNODA documentation

Contact

Gaynor Kim-Sing

Global Stewardship Council

gstewardship2026@proton.me

This document is published under CC0 (Public Domain). States Parties, civil society organizations, and individuals may reproduce, quote, translate, and distribute it without restriction.
